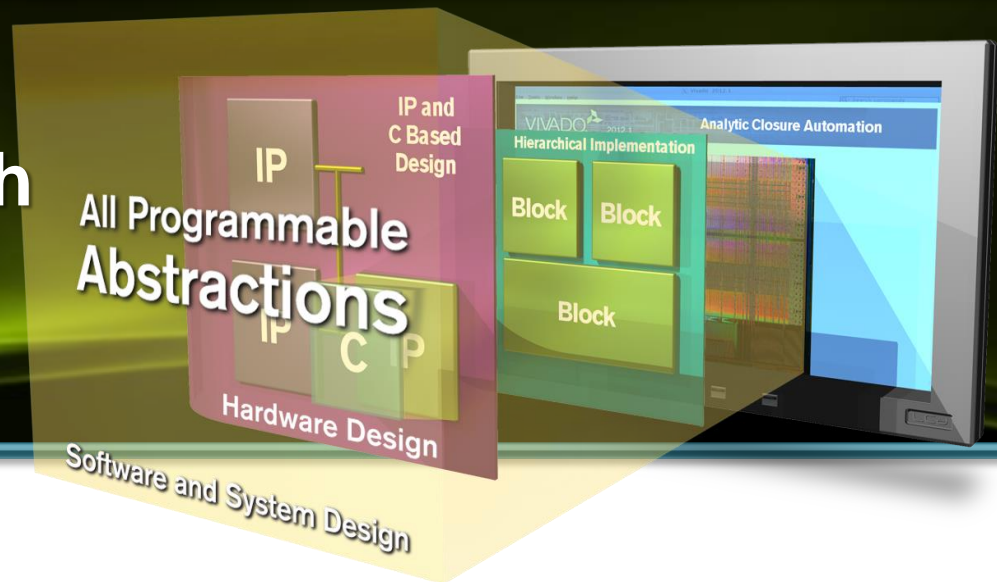


# Getting started with with Vivado HLS: FIR filtering



**Daniele Bagni**  
**DSP Specialist for EMEA**  
*daniele.bagni@xilinx.com*

# Outlines

- **HLS concepts: Initialization Interval and Latency**
- **FIR filter design: the C++ code**
- **FIR filter design: optimization directives and related performance**
  - Solution0: baseline
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
  - Solution3: directive LATENCY to limit loop iteration latency between min and max clock cycles
- **FIR filter design: performance summary**

# Outlines

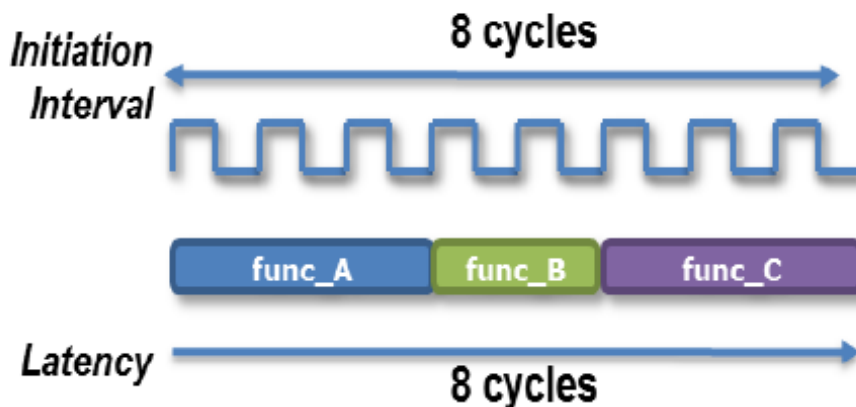
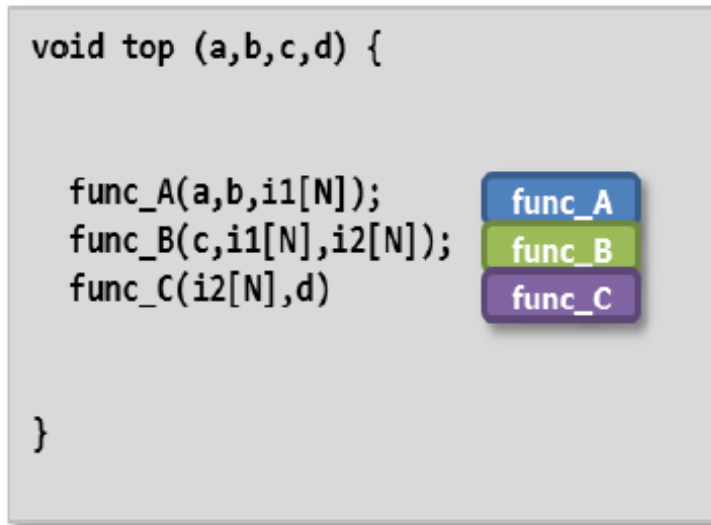
- **HLS concepts: Initialization Interval and Latency**
- FIR filter design: the C++ code
- FIR filter design: optimization directives and related performance
  - Solution0: baseline
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
  - Solution3: directive LATENCY to limit loop iteration latency between min and max clock cycles
- FIR filter design: performance summary

# Rationale

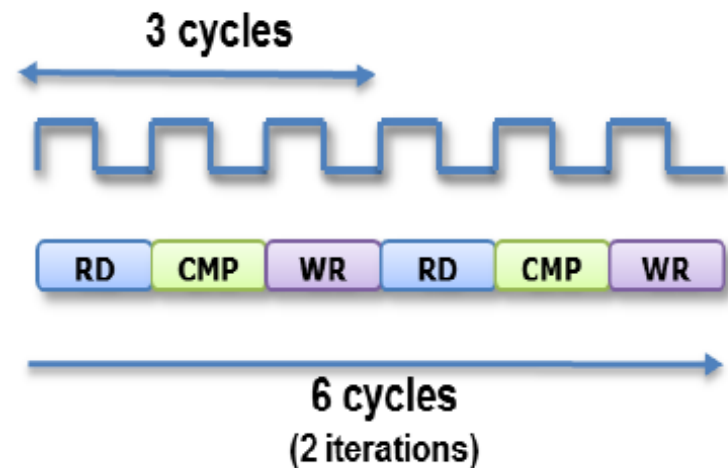
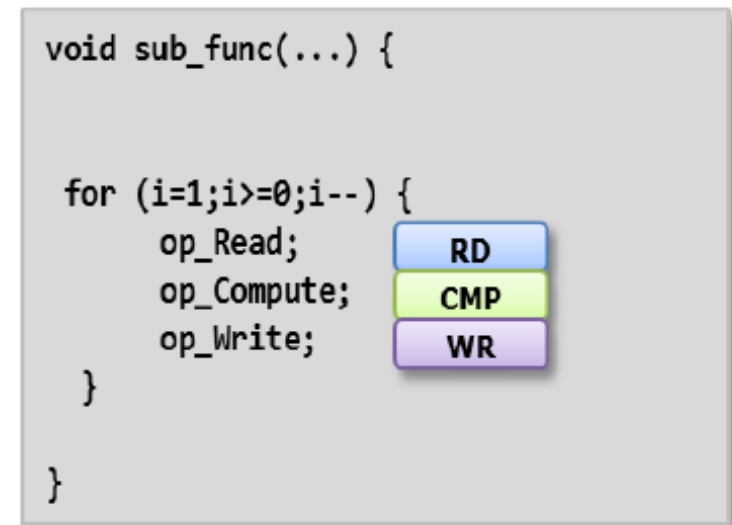
- Using a simple N=16 taps FIR filter modelled in C++ to illustrate few High Level Synthesis (HLS) directives that affect the effective output data rate by changing the Initialization Interval and therefore the Latency.
- “The **Initiation Interval**, often called the Interval or the **II**, is the number of clock cycles between when the task can start to accept new input data” (from UG902)
- “**Latency** is the number of clock cycles it takes to produce an output value” (from UG902)

# Initialization Interval and Latency 1/6

## Function Level



## Loop level



# Initialization Interval and Latency 2/6

## ➤ Latency:

- In the function example the latency of the design is 8 clock cycles: it takes 8 clock cycles from the start of function top until and output can be written.
- In the loop example, it takes 3 clock cycles to execute each iteration of the loop and a total latency of 6 clock cycles to execute all loop iterations.

## ➤ Initialization Interval:

- In the function example, func A cannot be executed again until func C is finished. Because it takes 8 clock cycles until func C completes, the II of function top is 8 clock cycles: it runs 8 clock cycles before it can start processing new input data.
- The loop example takes 6 clock cycles to execute all loop transactions, however it can accept a new input every 3 clock cycle and so the loop II is 3.

# Initialization Interval and Latency 3/6

- In both these examples, everything operates sequentially, since there is no concurrency, the latency and II are the same for both tasks.
- **Pipelining optimization** enables a sequential C description to be implemented as a concurrent hardware implementation.
- Next Figure shows the result when High-Level Synthesis is used to pipeline the sub-functions inside function top and the operations in the loop.
- With pipelining, both tasks can achieve higher performance. At the function level, **dataflow optimization** allows the sub-functions (A, B and C ) to execute as soon as data is available.

# Initialization Interval and Latency 4/6

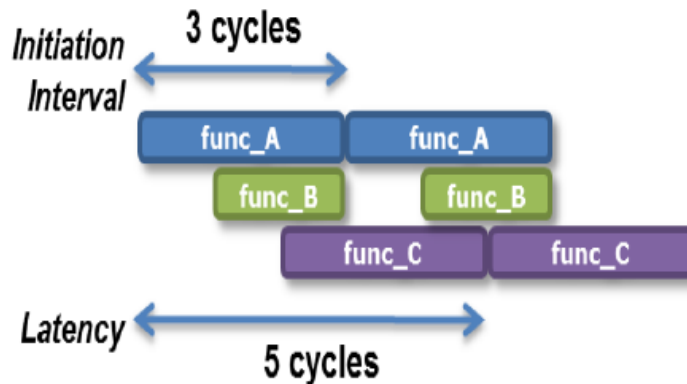
## Function Level

```
void top (a,b,c,d) {
```

```
    func_A(a,b,i1[N]);  
    func_B(c,i1[N],i2[N]);  
    func_C(i2[N],d)
```

func\_A  
func\_B  
func\_C

```
}
```



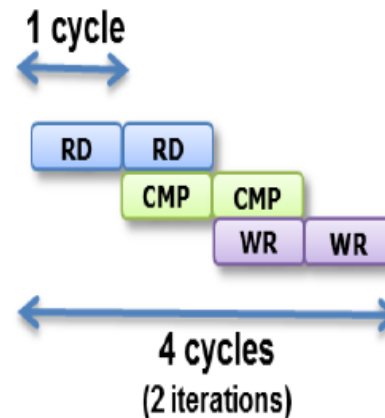
## Loop level

```
void sub_func(...) {
```

```
    for (i=1;i>=0;i--) {  
        op_Read;  
        op_Compute;  
        op_Write;  
    }
```

RD  
CMP  
WR

```
}
```





# Initialization Interval and Latency 5/6

- If function A starts to produce data before it has completely finished, func B can start accepting that data as soon as it is ready. It does not have to wait for func A to complete. Similarly, func C can start to execute as soon as data becomes available from func B. More importantly, func A can start the next transaction before func C has completed the current transaction.
- For the function top original design the latency and II were both 8 clock cycles. In this pipelined version of the same design, the latency is only 5 clock cycles and the II is 3 clock cycles: the pipelined design outputs data in less time and can accept data at a faster rate (almost 3X).

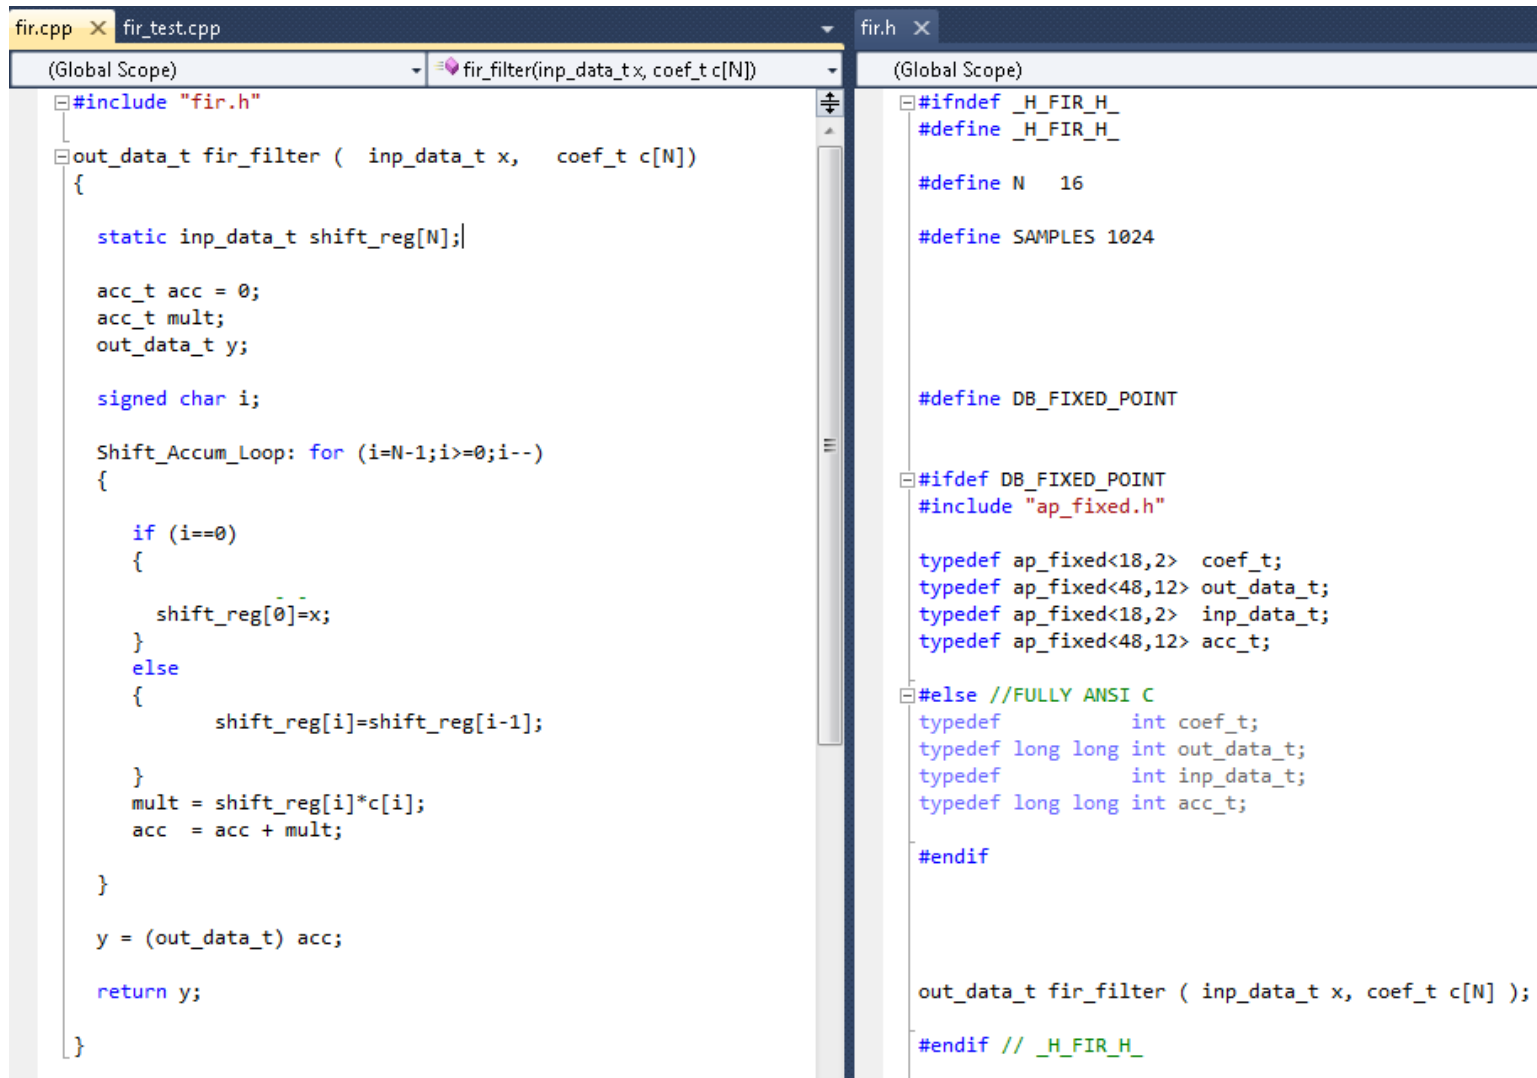
# Initialization Interval and Latency 6/6

- Performance improvements can be seen also for a pipelined implementation of the loop example. As soon as the read operation has completed, the next read operation can start.
- The performance of the loop can improve from an  $II=3$  to an  $II=1$  and by performing overlapped loop iterations, the overall loop latency can be reduced to 4 (even though the latency of each iteration remains at 3).

# Outlines

- HLS concepts: Initialization Interval and Latency
- **FIR filter design: the C++ code**
- FIR filter design: optimization directives and related performance
  - Solution0: baseline
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
  - Solution3: directive LATENCY to limit loop iteration latency between min and max clock cycles
- FIR filter design: performance summary

# FIR filter design: fir.cpp and fir.h design files



```

fir.cpp
(GLOBAL SCOPE)
#include "fir.h"
out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
{
    static inp_data_t shift_reg[N];

    acc_t acc = 0;
    acc_t mult;
    out_data_t y;

    signed char i;

    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
    {
        if (i==0)
        {
            shift_reg[0]=x;
        }
        else
        {
            shift_reg[i]=shift_reg[i-1];
        }
        mult = shift_reg[i]*c[i];
        acc = acc + mult;
    }

    y = (out_data_t) acc;

    return y;
}

fir.h
(GLOBAL SCOPE)
#ifndef _H_FIR_H_
#define _H_FIR_H_

#define N 16

#define SAMPLES 1024

#define DB_FIXED_POINT

#ifdef DB_FIXED_POINT
#include "ap_fixed.h"

typedef ap_fixed<18,2>  coef_t;
typedef ap_fixed<48,12> out_data_t;
typedef ap_fixed<18,2>  inp_data_t;
typedef ap_fixed<48,12> acc_t;
#else //FULLY ANSI C
typedef int coef_t;
typedef long long int out_data_t;
typedef int inp_data_t;
typedef long long int acc_t;
#endif

out_data_t fir_filter ( inp_data_t x, coef_t c[N] );

#endif // _H_FIR_H_

```

# FIR filter design: fir\_test.cpp testbench file

```
#include "fir.h"

int main (void)
{
    inp_data_t signal[SAMPLES];
    out_data_t output[SAMPLES], reference[SAMPLES];
    coef_t taps[N];

    FILE *fp1, *fp2, *fp3;
    int i, ret_value;
    float val1, val2;
    float diff, tot_diff;
    int val3;

    tot_diff = 0;

    // LOAD FILTER COEFFICIENTS
    fp1=fopen("./data/fir_coeff.dat","r");
    for (i=0;i<N;i++)
    {
        fscanf(fp1,"%f\n", &val1);
        taps[i] = (coef_t) val1;
        fprintf(stdout,"taps[%4d]=%10.5f\n", i, taps[i].to_double());
    }
    fclose(fp1);

    // LOAD INPUT DATA AND REFERENCE RESULTS
    fp2=fopen("./data/input.dat","r");
    fp3=fopen("./data/ref_res.dat","r");

    for (i=0;i<SAMPLES;i++)
    {
        fscanf(fp2,"%f\n", &val1);
        signal[i] = (inp_data_t) val1;
        fscanf(fp3,"%f\n", &val2);
        reference[i] = (out_data_t) val2;
    }
    fclose(fp2);
    fclose(fp3);

    fclose(fp2);
    fclose(fp3);

    // CALL DESIGN UNDER TEST
    for (i=0;i<SAMPLES;i++)
    {
        output[i] = fir_filter(signal[i], taps);
    }

    // WRITE OUTPUT RESULTS
    fp1=fopen("./data/out_res.dat","w");
    for (i=0;i<SAMPLES;i++)
    {
        fprintf(fp1,"%10.5f\n", output[i].to_double());
    }
    fclose(fp1);

    // CHECK RESULTS
    for (i=0;i<SAMPLES;i++)
    {
        diff = output[i].to_double() - reference[i].to_double();
        if (i<64) fprintf(stdout,"output[%4d]=%10.5f \t reference[%4d]=%10.5f\n",
            i, output[i].to_double(), i, reference[i].to_double() );

        diff = fabs(diff);
        tot_diff += diff;
    }
    fprintf(stdout, "TOTAL ERROR =%f\n",tot_diff);

    if (tot_diff < 1.0)
    {
        fprintf(stdout, "\nTEST PASSED!\n");
        ret_value = 0;
    }
    else
    {
        fprintf(stdout, "\nTEST FAILED!\n");
        ret_value = 1;
    }

    return ret_value;
}
```

# Outlines

- HLS concepts: Initialization Interval and Latency
- FIR filter design: the C++ code
- **FIR filter design: optimization directives and related performance**
  - Solution0: baseline
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
  - Solution3: directive LATENCY to limit loop iteration latency between min and max clock cycles
- FIR filter design: performance summary

# HLS solution0: first trials

- solution0\_a: baseline
- solution0\_b: Loop pipelining
- solution\_c: as solution\_b + partitioning the shift\_reg array
- solution\_d: as solution\_c + loop unrolling

## Performance Estimates

### Timing (ns)

Clock		solution0	solution0_b	solution0_c	solution0_d	solution0_e
default	Target	10.00	10.00	10.00	10.00	10.00
	Estimated	6.38	6.38	8.62	8.28	8.28

### Latency (clock cycles)

		solution0	solution0_b	solution0_c	solution0_d	solution0_e
Latency	min	97	37	20	11	2
	max	97	37	20	11	2
Interval	min	98	38	21	12	3
	max	98	38	21	12	3

## Utilization Estimates

	solution0	solution0_b	solution0_c	solution0_d	solution0_e
BRAM_18K	0	0	0	0	0
DSP48E	1	1	1	16	16
FF	202	171	466	1084	927
LUT	133	142	269	648	627

# Outlines

- HLS concepts: Initialization Interval and Latency
- FIR filter design: the C++ code
- **FIR filter design: optimization directives and related performance**
  - Solution0: directive LATENCY to limit loop iteration latency between min and max clock cycles
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
- FIR filter design: performance summary



# HLS solution1: II=1 at top level

- **solution1\_II1:**
  - II=1 at top level
- **solution1\_II2:**
  - II=2 at top level + allocation mul=8
- **solution1\_II4:**
  - II=4 at top level + allocation mul=4
- **solution1\_II8:**
  - II=8 at top level + allocation mul=2
- **solution1\_II16:**
  - II=16 at top level + allocation mul=1

## Performance Estimates

### Timing (ns)

Clock		solution1_II1	solution1_II2	solution1_II4	solution1_II8	solution1_II16
default	Target	10.00	10.00	-	10.00	10.00
	Estimated	8.28	8.28	-	8.28	8.28
ap_clk	Target	-	-	10.00	-	-
	Estimated	-	-	8.28	-	-

### Latency (clock cycles)

		solution1_II1	solution1_II2	solution1_II4	solution1_II8	solution1_II16
Latency	min	2	3	5	9	17
	max	2	3	5	9	17
Interval	min	1	2	4	8	16
	max	1	2	4	8	16

## Utilization Estimates

	solution1_II1	solution1_II2	solution1_II4	solution1_II8	solution1_II16
BRAM_18K	0	0	0	0	0
DSP48E	16	8	4	2	1
FF	927	640	610	654	662
LUT	626	916	773	776	823

# HLS solution1\_II1: directives to achieve II=1

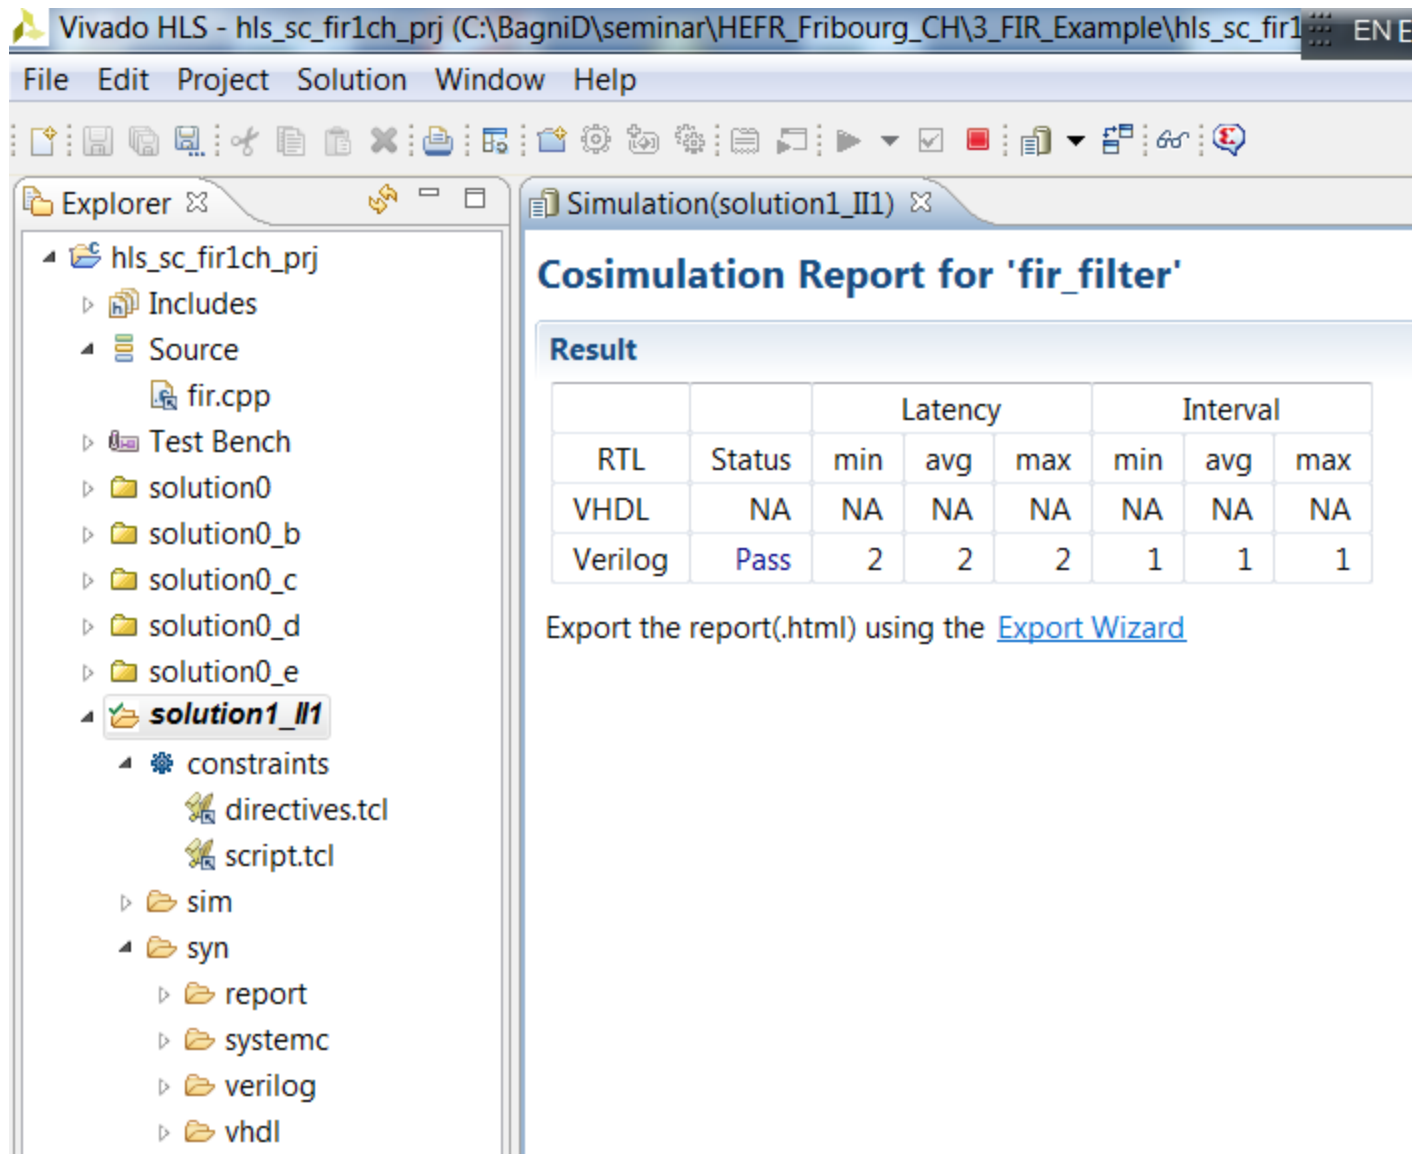
The screenshot displays the Vivado HLS IDE interface. The top menu bar includes File, Edit, Project, Solution, Window, and Help. The toolbar contains icons for various development actions. The left pane shows the Explorer view with the project structure for 'hls\_sc\_fir1ch\_prj', including folders for Includes, Source, Test Bench, and solution directories. The main editor window shows the 'fir.cpp' file with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34 }
35
36 }
37
```

The right pane shows the Directive view for the 'fir filter' block. The directives are as follows:

- % HLS PIPELINE II=1 (highlighted with a blue box)
- \*1 shift\_reg
- % HLS ARRAY\_PARTITION partition variable=shift\_reg complete dim=1
- x
- % HLS INTERFACE ap\_none register port=x
- c
- % HLS ARRAY\_PARTITION partition variable=c complete dim=1
- Shift\_Accum\_Loop

# HLS solution1\_II1: RTL/C Cosimulation report



The screenshot shows the Vivado HLS interface for project 'hls\_sc\_fir1ch\_prj'. The Explorer pane on the left displays the project structure, with 'solution1\_II1' selected. The Simulation pane on the right shows the 'Cosimulation Report for \'fir\_filter\''. The report includes a table with the following data:

RTL	Status	Latency			Interval		
		min	avg	max	min	avg	max
VHDL	NA	NA	NA	NA	NA	NA	NA
Verilog	Pass	2	2	2	1	1	1

Below the table, it states: 'Export the report(.html) using the [Export Wizard](#)'.

# HLS solution1\_II1: Export IP

## ➤ Results after Place-And-Route

### Export Report for 'fir\_filter'

#### General Information

Report date: Sun Jul 19 17:31:50 +0200 2015  
Device target: xc7z020clg484-1  
Implementation tool: Xilinx Vivado v.2015.1

#### Resource Usage

	Verilog
SLICE	156
LUT	566
FF	98
DSP	16
BRAM	0
SRL	10

#### Final Timing

	Verilog
CP required	10.000
CP achieved	9.580

# HLS solution1\_II2: directives to achieve II=2

The screenshot displays the Xilinx IDE interface for a project named 'hls\_sc\_fir1ch\_prj'. The main editor shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34
35
36 }
```

The right-hand pane shows the 'Directive' window for the 'fir\_filter' function. It lists the following HLS directives:

- HLS PIPELINE II=2** (highlighted with a blue box)
- HLS ALLOCATION instances=mul limit=8 operation**
- shift\_reg** (array partitioning)
- x** (variable)
- c** (variable)
- Shift\_Accum\_Loop** (loop unrolling)

The left-hand pane shows the project explorer with the following structure:

- hls\_sc\_fir1ch\_prj
  - Includes
  - Source
    - fir.cpp
  - Test Bench
    - fir\_test.cpp
  - data
  - solution0
  - solution0\_lat10
  - solution0\_lat20
  - solution1\_II1
  - solution1\_II16
  - solution1\_II2**
    - constraints
    - directives.tcl
    - script.tcl
    - syn
      - report
      - systemc
      - verilog
      - vhdl
  - solution1\_II4
  - solution1\_II8
  - solution1\_trial
  - solution2\_full\_pipeline
  - solution2\_unroll16
  - solution2\_unroll2
  - solution2\_unroll4
  - solution2\_unroll8

# HLS solution1\_II4: directives to achieve II=4

The screenshot displays the Vivado HLS IDE interface for a project named 'hls\_sc\_fir1ch\_prj'. The main editor shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34
35
36 }
```

The 'Directive' pane on the right lists the following directives for the 'fir\_filter' function:

- `% HLS PIPELINE II=4`
- `% HLS ALLOCATION instances=mul limit=4 operation`
- `*[] shift_reg`
- `% HLS ARRAY_PARTITION partition variable=shift_reg complete dim=1`
- `x`
- `% HLS INTERFACE ap_none register port=x`
- `c`
- `% HLS ARRAY_PARTITION partition variable=c complete dim=1`
- `*[] Shift_Accum_Loop`

The Explorer pane on the left shows the project structure, with 'solution1\_II4' selected under the 'solution0' directory.

# HLS solution1\_II8: directives to achieve II=8

The screenshot displays the Vivado HLS IDE interface for a project named 'hls\_sc\_fir1ch\_prj'. The main editor shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x, coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34
35
36 }
```

The 'Directive' pane on the right shows the following directives for the 'fir\_filter' function:

- % HLS ALLOCATION instances=mul limit=2 operation
- % HLS PIPELINE II=8
- shift\_reg
- % HLS ARRAY\_PARTITION partition variable=shift\_reg complete dim=1
- x
- % HLS INTERFACE ap\_none register port=x
- c
- % HLS ARRAY\_PARTITION partition variable=c complete dim=1
- Shift\_Accum\_Loop

The Explorer pane on the left shows the project structure, including the 'solution1\_II8' directory which contains 'constraints', 'directives.tcl', 'script.tcl', and 'syn'.

# HLS solution1\_II16: directives to achieve II=16

The screenshot displays the Vivado HLS environment for a project named 'hls\_sc\_fir1ch\_prj'. The Project Explorer on the left shows the file structure, with 'solution1\_II16' selected. The Code Editor in the center shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x, coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34 }
35
36 }
```

The Outline/Directive pane on the right shows the directives for the 'fir\_filter' function, with a blue box highlighting the first two directives:

- % HLS ALLOCATION instances=mul limit=1 operation
- % HLS PIPELINE II=16

Other directives visible include:

- \*[] shift\_reg
- % HLS ARRAY\_PARTITION partition variable=shift\_reg complete dim=1
- x
- % HLS INTERFACE ap\_none register port=x
- c
- % HLS ARRAY\_PARTITION partition variable=c complete dim=1
- Shift\_Accum\_Loop



# Outlines

- HLS concepts: Initialization Interval and Latency
- FIR filter design: the C++ code
- **FIR filter design: optimization directives and related performance**
  - Solution0: directive LATENCY to limit loop iteration latency between min and max clock cycles
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
- FIR filter design: performance summary

# HLS solution2: loop unrolling

- **solution2\_unr2:**
  - Loop unrolling by 2
- **solution2\_unr4:**
  - Loop unrolling by 4
- **solution2\_unr8:**
  - Loop unrolling by 8
- **solution2\_unr16:**
  - Loop unrolling by 16
- **solution2\_unr16\_fullpip:**
  - Loop unrolling by 16 plus pipeline at top level

## Performance Estimates

### Timing (ns)

Clock		solution2_unroll2	solution2_unroll4	solution2_unroll8	solution2_unroll16	solution2_full_pipeline
default	Target	10.00	10.00	10.00	10.00	10.00
	Estimated	8.62	8.62	8.62	8.28	8.28

### Latency (clock cycles)

		solution2_unroll2	solution2_unroll4	solution2_unroll8	solution2_unroll16	solution2_full_pipeline
Latency	min	33	25	23	2	2
	max	33	25	23	2	2
Interval	min	34	26	24	3	1
	max	34	26	24	3	1

## Utilization Estimates

	solution2_unroll2	solution2_unroll4	solution2_unroll8	solution2_unroll16	solution2_full_pipeline
BRAM_18K	0	0	0	0	0
DSP48E	2	4	8	16	16
FF	537	730	1120	927	927
LUT	983	1443	2652	627	626

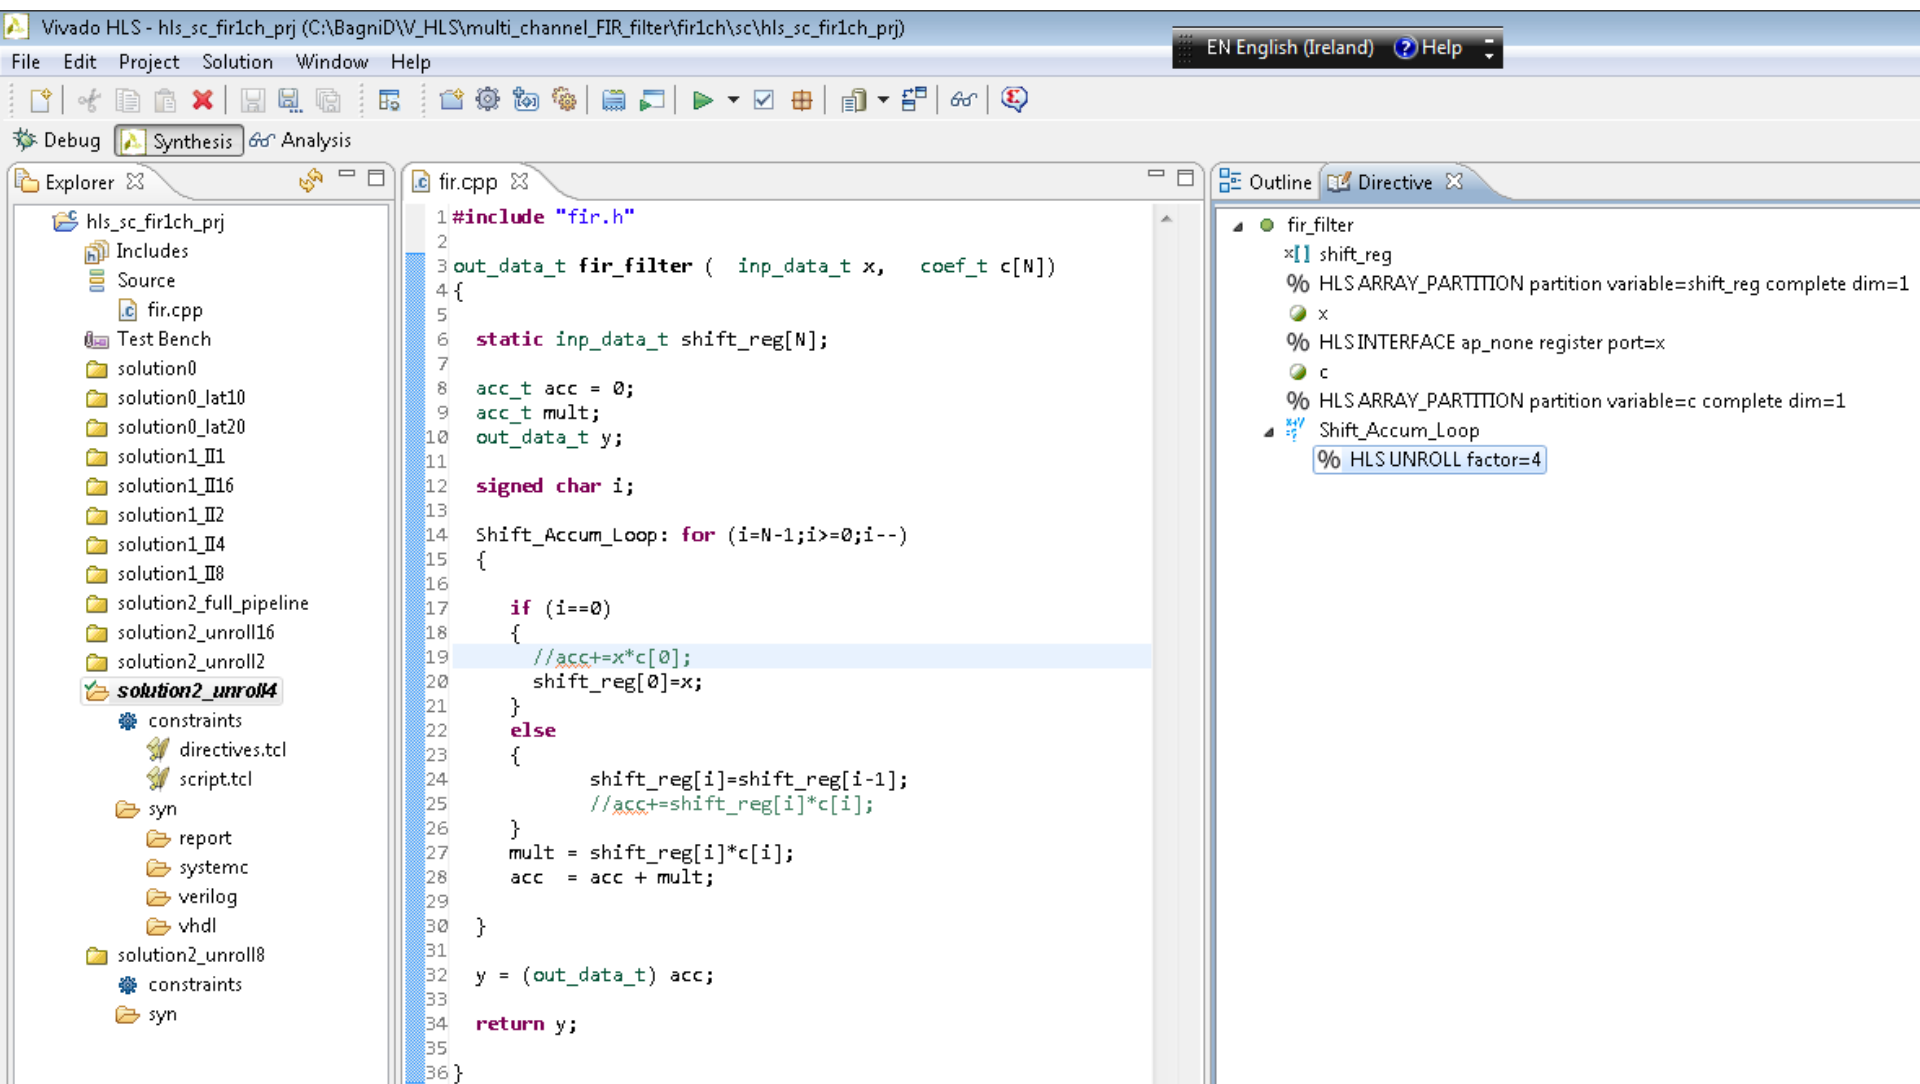
# HLS solution2\_unr2: loop unrolling by 2

The screenshot displays the Vivado HLS interface for a project named 'hls\_sc\_fir1ch\_prj'. The main editor shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34
35
36 }
```

The Explorer panel on the left shows the project structure, with 'solution2\_unroll2' selected. The Outline panel on the right shows the hierarchy of the 'fir\_filter' function, highlighting the 'Shift\_Accum\_Loop' and its 'HLS UNROLL factor=2'.

# HLS solution2\_unr4: loop unrolling by 4



# HLS solution2\_unr8: loop unrolling by 8

The screenshot displays the Vivado HLS interface for a project named 'hls\_sc\_fir1ch\_prj'. The Explorer pane on the left shows the project structure, including 'Source' files and various 'solution' folders. The 'fir.cpp' file is selected in the Source pane. The main editor window shows the C++ code for the 'fir\_filter' function, which implements a FIR filter with loop unrolling by 8. The code is as follows:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34
35
36 }
```

The Outline pane on the right shows the hierarchy of the code, including the 'fir\_filter' function and the 'Shift\_Accum\_Loop' loop. The 'Shift\_Accum\_Loop' is highlighted, showing the unrolling factor of 8.

# HLS solution2\_unr16: loop unrolling by 16

The screenshot displays the Vivado HLS IDE interface for a project named 'hls\_sc\_fir1ch\_prj'. The main editor shows the C++ source file 'fir.cpp' with the following code:

```
1 #include "fir.h"
2
3 out_data_t fir_filter ( inp_data_t x,  coef_t c[N])
4 {
5
6     static inp_data_t shift_reg[N];
7
8     acc_t acc = 0;
9     acc_t mult;
10    out_data_t y;
11
12    signed char i;
13
14    Shift_Accum_Loop: for (i=N-1;i>=0;i--)
15    {
16
17        if (i==0)
18        {
19            //acc+=x*c[0];
20            shift_reg[0]=x;
21        }
22        else
23        {
24            shift_reg[i]=shift_reg[i-1];
25            //acc+=shift_reg[i]*c[i];
26        }
27        mult = shift_reg[i]*c[i];
28        acc = acc + mult;
29    }
30
31    y = (out_data_t) acc;
32
33    return y;
34 }
35
36
37
```

The 'Shift\_Accum\_Loop' is highlighted in blue. The 'Outline' pane on the right shows the project structure, including the 'fir\_filter' function and the 'Shift\_Accum\_Loop' loop, which is annotated with '% HLS UNROLL factor=16'.

The 'Explorer' pane on the left shows the project hierarchy, including the 'solution2\_unroll16' directory, which contains the 'constraints', 'directives.tcl', 'script.tcl', 'syn', 'report', 'systemc', 'verilog', and 'vhdl' subdirectories.

# Outlines

- HLS concepts: Initialization Interval and Latency
- FIR filter design: the C++ code
- **FIR filter design: optimization directives and related performance**
  - Solution0: baseline
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
  - Solution3: directive LATENCY to limit loop iteration latency between min and max clock cycles
- FIR filter design: performance summary

# HLS solution3: latency

## ➤ solution3:

- Partitioning arrays coeff and shift\_reg

## ➤ Solution3\_lat10:

- as solution3 + latency=10

## ➤ Solution3\_lat20:

- as solution3 + latency=20

### Performance Estimates

#### ▣ Timing (ns)

Clock		solution3	solution3_lat10	solution3_lat20
default	Target	10.00	10.00	10.00
	Estimated	8.62	8.62	8.62

#### ▣ Latency (clock cycles)

		solution3	solution3_lat10	solution3_lat20
Latency	min	49	161	321
	max	49	177	337
Interval	min	50	162	322
	max	50	178	338

### Utilization Estimates

	solution3	solution3_lat10	solution3_lat20
BRAM_18K	0	0	0
DSP48E	1	1	1
FF	418	492	502
LUT	373	342	354



# Outlines

- HLS concepts: Initialization Interval and Latency
- FIR filter design: the C++ code
- FIR filter design: optimization directives and related performance
  - Solution0: directive LATENCY to limit loop iteration latency between min and max clock cycles
  - Solution1: directives PIPELINE II=x and ALLOCATION (this last one to limit the amount of mult operations)
  - Solution2: directive UNROLL to unroll the loop by x times
- **FIR filter design: performance summary**

# FIR Filter design: performance summary

solution name	clock freq MHZ	Latency	Interval	data rate MSPS	DSP48	FF	LUT
solution0_a: baseline	100	97	98	1.02	1	202	133
solution0_b: Loop pipeline	100	37	38	2.63	1	171	142
solution0_c: +partitioning shiftreg	100	20	21	4.76	1	466	269
solution0_d: +Loop unrolling	100	11	12	8.33	16	1084	648
solution0_e: +partitionin coeff array	100	2	3	33.33	16	927	627
solution1_II1	100	2	1	100.00	16	927	626
solution1_II2	100	3	2	50.00	8	640	916
solution1_II4	100	5	4	25.00	4	610	773
solution1_II8	100	9	8	12.50	2	654	776
solution1_II16	100	17	16	6.25	1	662	823
solution2_unroll2	100	33	34	2.94	2	537	983
solution2_unroll4	100	25	26	3.85	4	730	1443
solution2_unroll8	100	23	24	4.17	8	1120	2652
solution2_unroll16	100	2	3	33.33	16	927	627
solution2_unroll_full_pipelined	100	2	1	100.00	16	927	626
solution3	100	49	50	2.00	1	418	373
solution3_latency10	100	177	178	0.56	1	492	342
solution3_latency20	100	337	338	0.30	1	502	354

# FIR Filter design: HLS directives summary

## ➤ So far we have used:

- ARRAY\_PARTITION
- ALLOCATION
- PIPELINE
- UNROLL
- LATENCY
- LOOP TRIP\_COUNT

# Agenda

- 1) FIR filter case: 1 channel, in fractional fixed-point precision
  - ARRAY\_PARTITION, ALLOCATION, PIPELINE, UNROLL
  - LATENCY, LOOP TRIP\_COUNT
- 2) FIR filter case: 1 channel, in Floating-Point
- 3) Floating-Point Accumulator
- 4) Dependency
- 5) image Histogram computation and equalization
- 6) Siemens' application "Gamma LUT" case study
- 7) Image Processing: from HLS to IPI, from IPI to ZC702 board
- 8) cordic arctan2
- 9) cordic sqrt